

Alessandro Pietro Serra

PhD Student

- ✉ alex.serra.pietro@gmail.com
- 📍 Trieste, Italy
- .linkedin www.linkedin.com/in/alessandro-pietro-serra
- 🔗 <https://alexsserra98.github.io/>
- GitHub github.com/alexsserra98

Education

PhD: Theoretical and Scientific Data Science, SISSA • <i>Research Interests:</i> Interpretability, Causality, AI Safety	10/2024 – Present Trieste, Italy
MSc: Data Science and Scientific Computing, <i>University of Trieste</i> • <i>Curriculum:</i> Foundations of AI and ML. • <i>Thesis:</i> "The representation landscape of few-shot learning and fine-tuning in large language models" @ NeurIPS2024 • <i>Final mark:</i> 110/110 with Honors	10/2021 – 03/2024 Trieste, Italy
BSc: Mathematics for Engineering, Polytechnic of Turin	10/2018 – 10/2021 Turin, Italy

Professional Experience

Research Intern, Area Science Park • Investigated the effects of <i>fine-tuning</i> and <i>in-context learning</i> on the geometry of activations in LLaMA-family models, leading to a NeurIPS 2024 publication • Characterized at multiple scales the geometry of internal representations, applying clustering and intrinsic dimensionality estimators for macroscopic structure and nearest-neighbor analysis for local behavior.	10/2023 – 10/2024 Trieste, Italy
Machine Learning Engineer Intern, Aindo • Adapted, trained, and deployed a transformer-based model for tabular synthetic data generation. • Extended the company's core library by implementing a survival analysis module to validate synthetic data quality.	03/2023 – 05/2023 Trieste, Italy
Data Analyst, Espereal Technologies • Developed and integrated Python-based data analysis and visualization tools for IoT data into the company's main web platform.	01/2021 – 03/2021 Turin, Italy

Publications

The Narrow Gate: Localized Image-Text Communication in Native Multimodal Models, Advances in Neural Information Processing Systems (NeurIPS 2025) Alessandro Serra*, Francesco Ortu*, Emanuele Panizon*, Lucrezia Valeriani, Lorenzo Basile, Alessio Ansuini, Diego Doimo, Alberto Cazzaniga	2025
The representation landscape of few-shot learning and fine-tuning in large language models., Advances in Neural Information Processing Systems (NeurIPS 2024) Diego Doimo, Alessandro Serra, Alessio Ansuini, Alberto Cazzaniga	2024

Skills

Python | C++ | C | R | SQL | Git | Bash | Docker | PyTorch | Pandas | MPI | OpenMP | Slurm

Awards

Scholarships "Training the Specialists of the Future",
MSc Data Science and Scientific Computing

01/03/2023

Workshops & Talks

Workshop: Interpretability in LLMs using Geometric and Statistical Methods

05/2025 | Amsterdam, Netherlands

Presenter, "*Localized Image-Text Communication in VLMs*" at University of Amsterdam